


Model post-training, machine learning engineering, research replication, and systems optimization.


PostTrainBench

Post-training language models under a fixed compute budget.




MLE-bench-30

Solving machine learning tasks from Kaggle competitions.



PaperBench

Replicating machine learning research from papers.




Training & Kernel Optimization

Improving model quality, training speed, and GPU kernel performance.


SOL-ExecBench

Optimizing GPU kernels toward hardware performance limits.



NanoChat AutoResearch

Improving language model quality under a fixed training budget.



NanoGPT SpeedRun

Reducing training time to target validation loss.

